

Software tools for Maximum Likelihood Estimation:MB

Lesson 2 - first RTMB & Derivatives

[Click here to view presentation online](#)

Jim Bence

1 December 2023


Outline:

- Demos and exercises of using nlminb and RTMB to estimate parameters
- All that derivative stuff
 - Definitions derivatives, partial derivatives, second derivatives, cross derivatives (aka mixed second derivatives), gradient vector and Hessian
- Methods to calculate/approximate derivatives
- Exercise: finite difference derivatives
- Using the Hessian and gradient vector in parameter searches and to assess uncertainty
- RTMB Nonlinear regression vonb example

Demos of estimating the mean the hard way

- Just using `nlminb`
- Using RTMB
- Exercise - change model to assume gamma rather than normal (in breakout groups)

What is a derivative



$$\frac{df(\theta)}{d\theta} = \lim_{h \rightarrow 0} \frac{f(\theta + h) - f(\theta)}{h}$$

Partial derivative

Function with multiple arguments but we treat all but one of them as constants, and calculate derivative with respect to just one! E.g.,


$$\frac{\partial f}{\partial \theta_2} = \lim_{h \rightarrow 0} \frac{f(\theta_1, \theta_2 + h, \theta_3) - f(\theta_1, \theta_2, \theta_3)}{h}$$

Second derivative

Just a derivative of a derivative


$$\frac{\partial^2 f}{\partial \theta^2} = \frac{\partial \frac{\partial f}{\partial \theta}}{\partial \theta}$$

Visualizing second derivatives



Concave function with negative second derivative

Visualizing second derivatives



Convex function with positive second derivative

Cross derivative (mixed second derivs)

$$\frac{\partial^2 f}{\partial \theta_1 \partial \theta_2} = \frac{\partial}{\partial \theta_2} \frac{\partial f}{\partial \theta_1}$$

An important and convenient fact:

$$\frac{\partial^2 f}{\partial \theta_1 \partial \theta_2} = \frac{\partial^2 f}{\partial \theta_2 \partial \theta_1}$$

Methods for calculating derivatives

- Analytical derivatives. Gold standard but not available for many complex models.
- Finite difference methods. Intuitive but slow and propagate errors.
- Automatic differentiation. Fast and accurate but requires specialized software.

Finite difference derivatives

Widely used, e.g., default of `nlinb` and Excel solver

Forward difference

$$\partial f / \partial \theta_i = \frac{f(\theta_i + h) - f(\theta_i)}{h}$$

h is semi-arbitrary but small relative to θ_i .

Central differences

$$\partial f / \partial \theta_i = \frac{f(\theta_i + h) - f(\theta_i - h)}{2h}$$

Exercise - finite difference derivatives

- For $g(x) = a + bX + \sin X$, use finite difference methods to calculate the derivative of $g(X)$ with respect to (wrt) X
 - for $x=1$, with $a=2$ and $b=0.5$. (answer approximately 1.040)
 - Repeat for $x=2$, $a=1$, and $b=1$. Answer approximately 0.5839.
- For $a=2$, $b=0.5$, $x=1$, and same function, use finite differences to find the second derivative wrt x (answer approximately -0.8415)

Automatic differentiation

- Uses repeated applications of chain rule:

$$\partial z / \partial \theta = [\partial z / \partial y][\partial y / \partial \theta]$$

- Simplest case. $y = f(\theta)$, $z = g(y)$, i.e., $z = g(f(\theta))$
- General case we care about:

$$NLL = f_1(f_2(f_3(\dots f_k(\theta)\dots)))$$

Gradient

Just a fancy term to mean the vector of derivatives of the NLL function with respect to each parameters (so if k parameters, then k elements)

$$g = \{\partial f / \partial \theta_1, \partial f / \partial \theta_2, \dots \partial f / \partial \theta_k\}^T$$

Hessian - a square symmetric matrix

$$H = \begin{bmatrix} \partial^2 f / \partial \theta_1^2 & \partial^2 f / \partial \theta_1 \partial \theta_2 & \dots & \partial^2 f / \partial \theta_1 \partial \theta_k \\ \partial^2 f / \partial \theta_2 \partial \theta_1 & \partial^2 f / \partial \theta_2^2 & \dots & \partial^2 f / \partial \theta_2 \partial \theta_k \\ \dots & \dots & \dots & \dots \\ \partial^2 f / \partial \theta_k \partial \theta_1 & \partial^2 f / \partial \theta_k \partial \theta_2 & \dots & \partial^2 f / \partial \theta_k^2 \end{bmatrix}$$

$$h_{i,j} = h_{j,i} = \frac{\partial^2 f}{\partial \theta_i \partial \theta_j} = \frac{\partial^2 f}{\partial \theta_j \partial \theta_i}$$

If the NLL were a quadratic function as it would be for linear normal model...

$$\theta_{\min} = \theta_{\text{start}} + H^{-1}g$$

where H^{-1} is the matrix inverse of H and $H^{-1}g$ is the product of the inverse of the Hessian and the gradient

Because our models generally not normal and linear, iterative searches...

1. specify starting values for parameters, $\underline{\theta}_0$
2. Replace $\underline{\theta}_0$ by $\underline{\theta}_1 = \underline{\theta}_0 + \delta_0$
3. Check gradient and Hessian and if at a minimum stop otherwise...
4. Return to step 2 but each time $\underline{\theta}_{i+1} = \underline{\theta}_i + \delta_i$

Newton step: $\underline{\delta}_i = H^{-1} \underline{g}$ evaluated at current params

Quasi-Newton method uses $\underline{\delta}_i = \lambda H^{-1} \underline{g}$ with Hessian approximated using search path, and λ a number less than 1

Using the Hessian to calculate asymptotic standard errors

- First some reminders
 - Parameter estimates are random variates that result from estimators (random variables)
 - The variance describes the variability of results from applying the estimation method, namely the expected squared deviation between an estimator and its expected value
 - What we report as a standard error for a parameter is the square-root of this variance.

The variance-covariance matrix


$$\Sigma = \begin{matrix} & \sigma_1^2 & \sigma_{1,2} & \dots & \sigma_{1,k} \\ \sigma_{2,1} & \sigma_2^2 & \dots & \sigma_{2,k} \\ \dots & \dots & \dots & \dots \\ \sigma_{k,1} & \sigma_{k,2} & \dots & \sigma_k^2 \end{matrix}$$

The asymptotic variance-covariance matrix

$$\hat{\Sigma} = H^{-1}$$


- Square-root of diagonal gives standard errors
- Off-diagonals are covariances
- The Hessian needs to be positive definite for the calculation
- If the Hessian is not positive definite its a problem!
- Delta method used to obtain SEs for derived quantities (using $\hat{\Sigma}$)

Musky vonB example



musky_vonb.dat

von Bertalanffy model

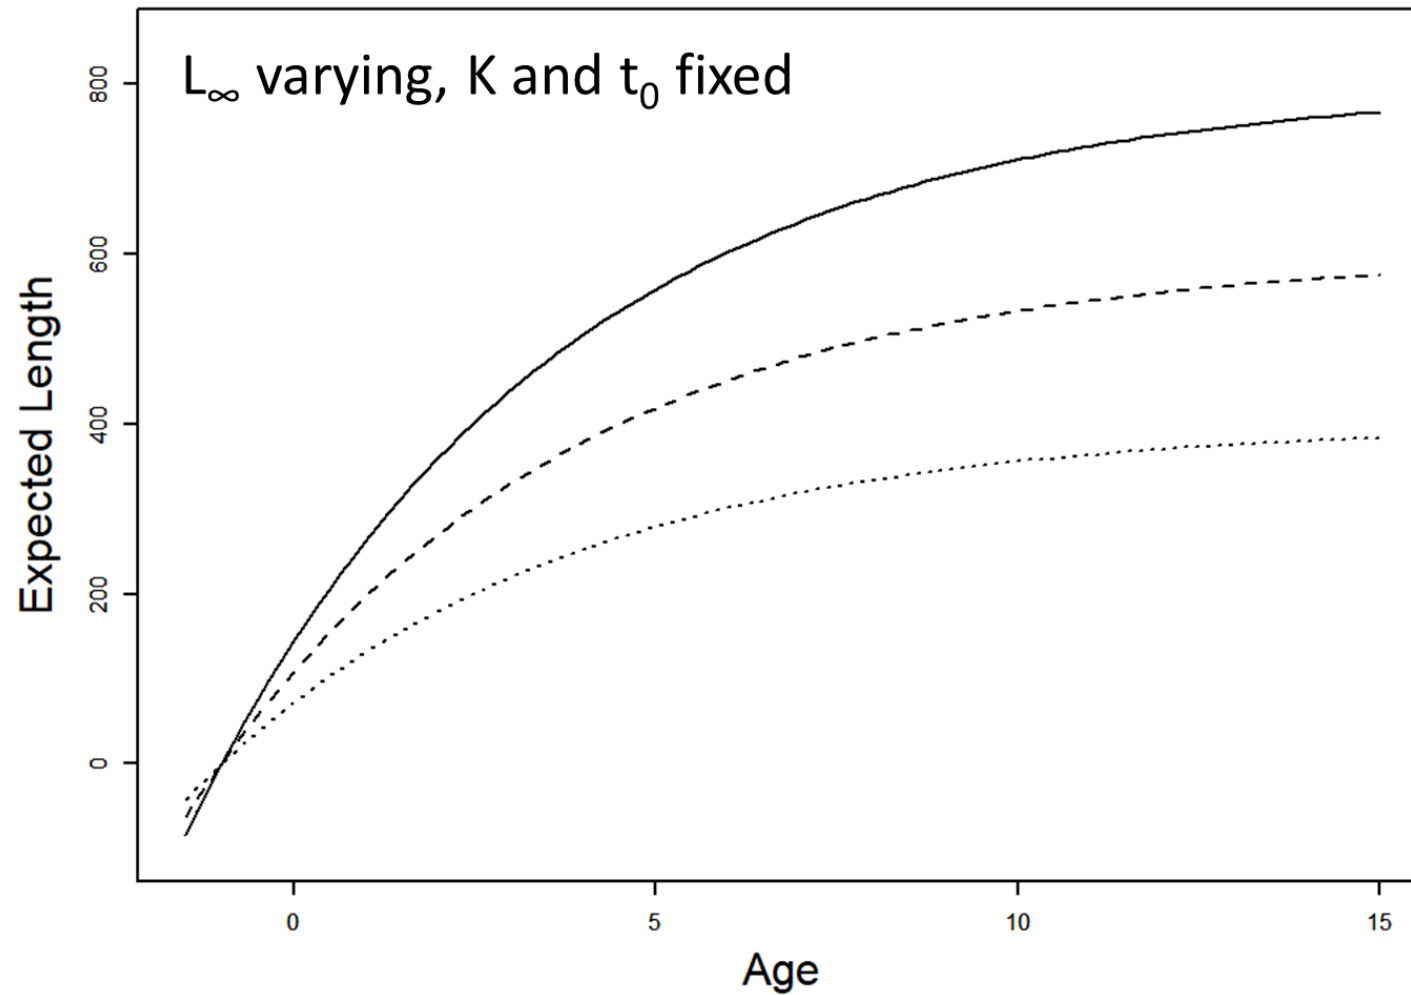


$$L_i = L_\infty \left(1 - e^{-K(a_i - t_0)}\right) + \varepsilon_i$$


$$\varepsilon_i \stackrel{iid}{\sim} N(0, \sigma^2)$$

$$L_i \sim N(L_{a_i}, \sigma^2)$$

Influence of L_{∞}



Influence of K



Musky vonB setup code

```
1 library(RTMB);  
2  
3 gmRdat = read.table("lesson2/data/musky_vonb.dat", head=T);  
4  
5 #Set up the data and starting value of parameters for RTMB  
6 gmdat = list(len_obs=gmRdat[, "Length"], age=gmRdat[, "Age"]);  
7 gmpar = list(log_linf=7, log_vbk=-1.6, t0=0, log_sd=4);
```

code for NLL for Musky vonb example

```
1  NLL_fun = function(par_lst){
2    getAll(gmdat,par_lst);
3    linf = exp(log_linf);
4    vbk = exp(log_vbk);
5    sd = exp(log_sd);
6    len_pred = linf * (1 - exp(-vbk * (age - t0)));
7    nll = -sum(dnorm(len_obs, len_pred, sd, TRUE));
8    atage_pred = linf * (1 - exp(-vbk * ((1:11) - t0)))
9    REPORT(atage_pred);
10   nll
11 }
```

Create model object and print predicted lengths before fitting model

```
1  obj <- MakeADFun(NLL_fun, gmpar);  
2  
3  GMreport=obj$report();  
4  GMreport
```

\$atage_pred

```
[1] 200.4870 364.3209 498.2026 607.6080 697.0118 770.0708 829.7731 878.5606  
[9] 918.4287 951.0082 977.6314
```

fit the model

```
1 fit = nlminb(obj$par, obj$fn, obj$gr);
```

```
outer mgc: 3572.344  
outer mgc: 115.0372  
outer mgc: 306.8376  
outer mgc: 101.0471  
outer mgc: 30.69436  
outer mgc: 135.1418  
outer mgc: 120.5931  
outer mgc: 25.99566  
outer mgc: 197.1401  
outer mgc: 122.5854  
outer mgc: 66.52428  
outer mgc: 36.54847  
outer mgc: 3.778352  
outer mgc: 19.59957  
outer mgc: 1.539833  
.
```

Get parameter uncertainties and convergence diagnostics

```
1 sdr = sdreport(obj);
```

```
outer mgc: 0.0001970265
outer mgc: 19.77567
outer mgc: 19.71682
outer mgc: 9.178336
outer mgc: 9.171702
outer mgc: 2.35869
outer mgc: 2.358392
outer mgc: 0.1198859
outer mgc: 0.1201142
```

```
1 sdr #summary(sdr)
```

```
sdreport(.) result
      Estimate Std. Error
log_linf  7.1488714 0.04083255
log_vbk  -1.2456407 0.16950408
t0        -0.7710237 0.35092437
log_sd     3.8876390 0.09128707
Maximum gradient component: 0.0001970265
```

Predicted lengths at age after model fitting

```
1 GMreport = obj$report();  
2 GMreport;
```

\$atage_pred

```
[1] 508.1491 699.3217 842.6904 950.2090 1030.8420 1091.3122 1136.6614  
[8] 1170.6709 1196.1760 1215.3035 1229.6480
```


vonB Exercises

- Change `REPORT(atage_pred)` to `ADREPORT(atage_pred)` and look at `sdreport` and summary of the `sdreport`
- Calculate a new variable equal to $\text{vbk} * \text{linf}$ as `ADREPORT`
- Change the model so data are assumed gamma distributed (with expected value given by vonB equation and constant variance)